

# Class 07: machine learning 1

AUTHOR


Lisa Chen (PID: A17082974)

##clustering methods The broad goal here is to find groupings (clusters) in your input data.

## Kmeans

1st let's make up some data to cluster

```
x <- rnorm(1000)
hist(x)
```



make a vector of length 60 with 30 points centered at -3 and 30 points centered at +3

```
tmp <-c(rnorm(30, mean=-3), rnorm(30, mean=3))
tmp
```

```

[1] -2.9518996 -0.7827734 -1.7318752 -3.1115964 -1.8610008
-2.2718214
[7] -3.9678016 -3.8656070 -3.9810432 -2.3832023 -2.6293429
-1.8678918
[13] -4.1068040 -3.4800152 -2.4794825 -3.6767563 -4.7963726
-3.4123336
[19] -3.2444365 -3.5443769 -3.0204870 -3.0085322 -2.9176768
-1.9424662
[25] -3.5753728 -4.1667347 -1.2990372 -1.5418209 -2.1225637
-4.2290965
[31]  3.7922418  4.9277223  2.9082018  2.5006931  3.6771648
3.7659838
[37]  1.4497013  3.1369309  2.9501651  4.0875572  3.6604628
2.6377046
[43]  3.6506235  2.8175113  2.9379289  3.5623454  3.0398269
3.7737301
[49]  3.4970648  3.7059298  3.3818532  1.2477756  3.6874507
2.1120629
[55]  2.1779770  3.6603735  0.9791801  2.8033344  1.7976011
2.7429565

```

create x and y data set with 2 groups of points.

```

x <- cbind(x=tmp, y=rev(tmp))
x

```


```

      x      y
[1,] -2.9518996  2.7429565
[2,] -0.7827734  1.7976011
[3,] -1.7318752  2.8033344
[4,] -3.1115964  0.9791801
[5,] -1.8610008  3.6603735
[6,] -2.2718214  2.1779770
[7,] -3.9678016  2.1120629
[8,] -3.8656070  3.6874507
[9,] -3.9810432  1.2477756
[10,] -2.3832023  3.3818532
[11,] -2.6293429  3.7059298
[12,] -1.8678918  3.4970648
[13,] -4.1068040  3.7737301
[14,] -3.4800152  3.0398269
[15,] -2.4794825  3.5623454
[16,] -3.6767563  2.9379289
[17,] -4.7963726  2.8175113

```

```
[18,] -3.4123336 3.6506235
[19,] -3.2444365 2.6377046
[20,] -3.5443769 3.6604628
[21,] -3.0204870 4.0875572
[22,] -3.0085322 2.9501651
[23,] -2.9176768 3.1369309
[24,] -1.9424662 1.4497013
[25,] -3.5753728 3.7659838
[26,] -4.1667347 3.6771648
[27,] -1.2990372 2.5006931
[28,] -1.5418209 2.9082018
[29,] -2.1225637 4.9277223
[30,] -4.2290965 3.7922418
[31,] 3.7922418 -4.2290965
[32,] 4.9277223 -2.1225637
[33,] 2.9082018 -1.5418209
[34,] 2.5006931 -1.2990372
[35,] 3.6771648 -4.1667347
[36,] 3.7659838 -3.5753728
[37,] 1.4497013 -1.9424662
[38,] 3.1369309 -2.9176768
[39,] 2.9501651 -3.0085322
[40,] 4.0875572 -3.0204870
[41,] 3.6604628 -3.5443769
[42,] 2.6377046 -3.2444365
[43,] 3.6506235 -3.4123336
[44,] 2.8175113 -4.7963726
[45,] 2.9379289 -3.6767563
[46,] 3.5623454 -2.4794825
[47,] 3.0398269 -3.4800152
[48,] 3.7737301 -4.1068040
[49,] 3.4970648 -1.8678918
[50,] 3.7059298 -2.6293429
[51,] 3.3818532 -2.3832023
[52,] 1.2477756 -3.9810432
[53,] 3.6874507 -3.8656070
[54,] 2.1120629 -3.9678016
[55,] 2.1779770 -2.2718214
[56,] 3.6603735 -1.8610008
[57,] 0.9791801 -3.1115964
[58,] 2.8033344 -1.7318752
[59,] 1.7976011 -0.7827734
[60,] 2.7429565 -2.9518996
```

```
plot(x)
```



```
k <- kmeans(x, centers=2)
k
```

K-means clustering with 2 clusters of sizes 30, 30

Cluster means:

	x	y
1	3.035669	-2.932341
2	-2.932341	3.035669

Clustering vector:

```
[1] 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2
2 1 1 1 1 1 1 1 1 1
[39] 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1
```

Within cluster sum of squares by cluster:

```
[1] 51.86166 51.86166
(between_SS / total_SS = 91.2 %)
```

Available components:

```
[1] "cluster"      "centers"      "totss"       "withinss"
"tot.withinss"
[6] "betweenss"    "size"         "iter"        "ifault"
```

Q. from your result object `k` how many points are in each cluster?

```
k$size
```

```
[1] 30 30
```

Q. what "component" of your result object details the cluster membership?

```
k$cluster
```

```
[1] 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2
2 1 1 1 1 1 1 1 1 1
[39] 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1
```


Q. cluster centers?

```
k$centers
```

```
      x      y
1  3.035669 -2.932341
2 -2.932341  3.035669
```


Q. Plot of our clustering results

```
plot(x, col=k$cluster)
points(k$centers, col="blue", pch=15, cex=2)
```



we can cluster into 4 groups

```
# kmeans  
k4 <- kmeans(x, centers=4)  
# plot results  
plot(x, col=k4$cluster)
```



A big limitation of kmeans is that it does what you ask even if you ask for silly clusters.

## Hierarchical Clustering

The main base R function for Hierarchical Clustering is `hclust()` unlike `kmeans()` you can not just pass it your data as input. You first need to calculate a distance matrix.


```
d <- dist(x)
hc <- hclust(d)
hc
```

Call:  
`hclust(d = d)`

Cluster method : complete  
Distance : euclidean  
Number of objects: 60

Use `plot()` to view results.

```
plot(hc)
abline(h=10, col="red")
```



To make the "cut" and get our cluster membership vector, we can use the `cutree()` function.


```
grps <- cutree(hc, h=10)
grps
```

```
[1] 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1
1 2 2 2 2 2 2 2 2 2
[39] 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2
```

Make a plot of our data colored by hclust results

```
plot(x, col=grps)
```





## Principal Component Analysis (PCA)

Here we will do Principal Component Analysis (PCA) on some food

```
url <- "https://tinyurl.com/UK-foods"
x <- read.csv(url, row.names = 1)
head(x)
```

	England	Wales	Scotland	N.Ireland
Cheese	105	103	103	66
Carcass_meat	245	227	242	267
Other_meat	685	803	750	586
Fish	147	160	122	93
Fats_and_oils	193	235	184	209
Sugars	156	175	147	139

```
#rownames(x) <- x[,1]
#x <- x[, -1]
#x
```

Q1. How many rows and columns are in your new data frame named

x? What R functions could you use to answer this questions?


```
dim(x)
```

```
[1] 17  4
```

Q2. Which approach to solving the 'row-names problem' mentioned above do you prefer and why? Is one approach more robust than another under certain circumstances?

Second approach because if you run the first code multiple times, it keeps eliminating the first column.

```
barplot(as.matrix(x), beside=T, col=rainbow(nrow(x)))
```




Q3: Changing what optional argument in the above barplot() function results in the following plot?

You can change the Beside Argument.

Q5: Generating all pairwise plots may help somewhat. Can you make

sense of the following code and resulting figure? What does it mean if a given point lies on the diagonal for a given plot?

```
pairs(x, col=rainbow(10), pch=16)
```



## PCA to the rescue

the main "base" R function for PCA is called `prcomp()`. Here we need to take the transpose of our input

```
pca <- prcomp( t(x) )
summary(pca)
```

Importance of components:

	PC1	PC2	PC3	PC4
Standard deviation	324.1502	212.7478	73.87622	2.921e-14
Proportion of Variance	0.6744	0.2905	0.03503	0.000e+00
Cumulative Proportion	0.6744	0.9650	1.00000	1.000e+00

Q. how much variance is captured in 2 PCs

96.5%

to make our main "PC score plot" (aka "PC1 vs PC2 plot", or "PC plot", or "ordination plot")

```
attributes(pca)
```

```
$names
```

```
[1] "sdev"      "rotation" "center"    "scale"     "x"
```

```
$class
```


```
[1] "prcomp"
```

we are after the `pca\$x` result component to make our main PCA plot

```
pca$x
```

	PC1	PC2	PC3	PC4
England	-144.99315	-2.532999	105.768945	-9.152022e-15
Wales	-240.52915	-224.646925	-56.475555	5.560040e-13
Scotland	-91.86934	286.081786	-44.415495	-6.638419e-13
N.Ireland	477.39164	-58.901862	-4.877895	1.329771e-13

```
mycols <-c("orange", "red", "blue", "darkgreen")
plot(pca$x[,1], pca$x[,2], col=mycols, pch=16, xlab="PC1 (67.4%
```



another important result is from PCA is how the original variables (in this case, the foods) contribute to the PCs.

this is contained in the `pca$rotation` object folks often call this the "loadings" or "contributions" to the PCs

```
pca$rotation
```

	PC1	PC2	PC3
PC4			
Cheese	-0.056955380	0.016012850	0.02394295
-0.409382587			
Carcass_meat	0.047927628	0.013915823	0.06367111
0.729481922			
Other_meat	-0.258916658	-0.015331138	-0.55384854
0.331001134			
Fish	-0.084414983	-0.050754947	0.03906481
0.022375878			
Fats_and_oils	-0.005193623	-0.095388656	-0.12522257
0.034512161			
Sugars	-0.037620983	-0.043021699	-0.03605745
0.024943337			
Fresh_potatoes	0.401402060	-0.715017078	-0.20668248
0.021396007			

Fresh_Veg	-0.151849942	-0.144900268	0.21382237
0.001606882			
Other_Veg	-0.243593729	-0.225450923	-0.05332841
0.031153231			
Processed_potatoes	-0.026886233	0.042850761	-0.07364902
-0.017379680			
Processed_Veg	-0.036488269	-0.045451802	0.05289191
0.021250980			
Fresh_fruit	-0.632640898	-0.177740743	0.40012865
0.227657348			
Cereals	-0.047702858	-0.212599678	-0.35884921
0.100043319			
Beverages	-0.026187756	-0.030560542	-0.04135860
-0.018382072			
Soft_drinks	0.232244140	0.555124311	-0.16942648
0.222319484			
Alcoholic_drinks	-0.463968168	0.113536523	-0.49858320
-0.273126013			
Confectionery	-0.029650201	0.005949921	-0.05232164
0.001890737			

we can make a plot along pc1

```
library(ggplot2)
contrib <- as.data.frame(pca$rotation)
ggplot(contrib) +
  aes(PC1, rownames(contrib)) +
  geom_col()
```

